

Biological Stats II : Lab 5

Gavin Fay

02/20/2025

Lab schedule

1/22: Introduction to R and R Studio, working with data

1/29: Visualizing data

2/05: Probability, linear modeling

2/12: Data wrangling, model summaries

2/20: Simulation, Resampling

2/26: Iteration

3/19: Creating functions, debugging

4/02: working with Spatial data

4/09: Flex: automated reporting & Quarto

Segue: Packages in R

Installing packages:

```
install.packages("packagename")
```

If in doubt, use the source:

```
install.packages("packagename", type="source")
```

Installing packages from github:

```
remotes::install_github("repositoryname")
```

Loading (attaching) packages into the workspace:

```
library(packagename)
```

Some people use `require()` instead of `library()`.

Don't do this!

`require()` is basically `try(library())`

Often description of a package and how to use its functions is in the form of a vignette.

`vignette()` lists available vignettes.

`vignette(packagename)` views the vignette for `packagename`.

Programming practices I

Use projects (RStudio)

Write scripts (or markdown files)

Include whitespace in code

- blank lines, spaces in functions

Use an editor with syntax highlighting

Use a style guide

Indent code

Use meaningful object names

Programming practices II

Test code

- Write smallest possible amount (e.g. 1 line).
- Knit early and often.
- Try simple examples that you know the answer to.
- Always assume that there will be an error somewhere.

View results / objects (e.g. with `print()`).

Plot results - are they what you expect?

Be careful when copying sections of code and changing a variable name (it's super common to forget to change all occurrences).

- hint: use your text editor's "Find: Replace all" functionality.

Commenting

R ignores everything on a line that follows a #

Comment at the top of your script.

- What the script does, your name, email, date started.

Comment before each function or section of code - What is the purpose of that section of code, what does it do?

- Comment the 'why' not the 'what'

Comment throughout:

- whenever an unusual function is used
- whenever the code is hard to understand
- whenever an algorithm is particularly useful

Commenting out code

When you make modifications to your code: - Copy the code that works then comment it out by prefixing it with #. - Change the new copy of the code.

If you need to revert to the old code, just remove the # before each line (“uncomment”).

ctrl+shift+C is a shortcut in Rstudio to comment/uncomment large blocks of code.

In .Rmd files, you can comment out blocks of the file using

```
<!--
```

```
Lines of text and code you want to not be included
```

```
-->
```


Permutation tests

```
library(moderndiver)
null_evals <- evals |>
  specify(score ~ gender) |>
  hypothesize(null = "independence") |>
  generate(reps = 1000, type = "permute") |>
  calculate(stat = "diff in means",
            order = c("male", "female"))
null_evals_ci <- null_evals |>
  summarize(
    l = quantile(stat, 0.025),
    u = quantile(stat, 0.975)
  )
score_means <- evals |> group_by(gender) |>
  summarize(avg_score = mean(score))
dscore <- score_means[2,2] - score_means[1,2]

percentile_ci <- null_evals |>
  get_confidence_interval(type = "percentile",
                        level = 0.95)
```



```
visualize(null_evals) +  
  shade_confidence_interval(endpoints = percentile_ci) +  
  geom_vline(xintercept = as.numeric(dscore), linetype = "dashed")
```



Lab exercise 1 - permutation test


Use the Laengelmavesi data in `../data/Laengelmavesi2.csv`

- a. Obtain the data for just the lengths of perch and bream.
- b. Plot the distribution of lengths for both species, and calculate the mean lengths for both species.
- c. Conduct a permutation test to assess whether the difference in mean length between bream and perch is statistically clear.
- d. Plot the distribution for the test statistic under the null hypothesis of no difference in length, and indicate the true value for the test statistic relative to the two-tailed 95th percent of the null hypothesis distribution.
- e. What are your conclusions?

Auto dataset

in 'ISLR' package

``geom_smooth()`` using formula = `'y ~ x'`



Bootstrapping

Recall:

- ▶ Have some data x_1, x_2, \dots, x_n and we are computing a statistic.
- ▶ Randomly draw n values from the data *with replacement* (same value can be drawn multiple times).
- ▶ Calculate statistic from new random pseudo-data.
- ▶ Repeat a large number of times to obtain the distribution:
 t_1, t_2, \dots, t_n .
- ▶ Resulting distribution is the bootstrap sampling distribution.
- ▶ Compute standard deviation and 95 percent confidence intervals from this. Finished.

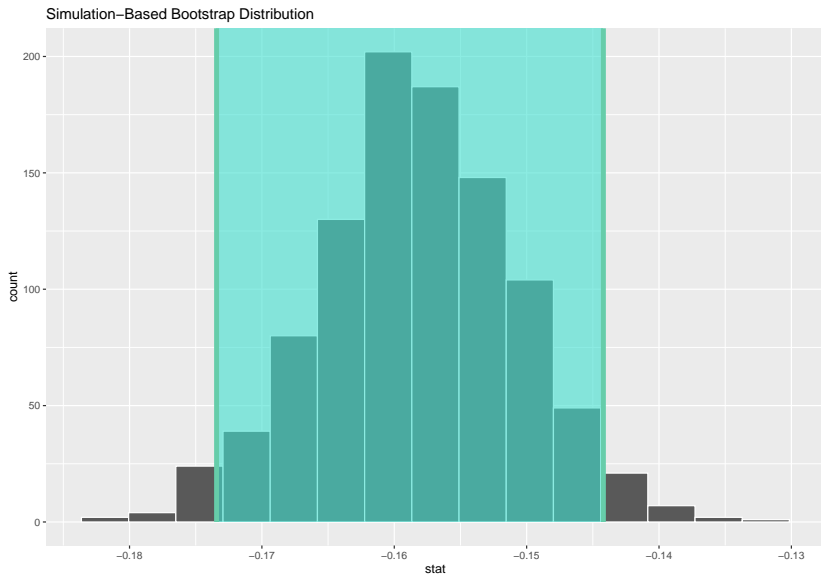
Performing generic case bootstrapping in R

Bootstrap estimates of slope from a linear model.

```
slope_bootstrap <- Auto |>
  specify(formula = mpg ~ horsepower) |>
  generate(reps = 1000, type = "bootstrap") |>
  calculate(stat = "slope")
percentile_ci <- slope_bootstrap |>
  get_confidence_interval(type = "percentile",
                          level = 0.95)

percentile_ci
# A tibble: 1 x 2
  lower_ci upper_ci
  <dbl>     <dbl>
1   -0.173   -0.144
```

```
visualize(slope_bootstrap) +  
  shade_confidence_interval(endpoints = percentile_ci)
```



Residual bootstrap

```
slope <- coef(lm(mpg~horsepower,data=Auto))[2]
```

```
lm1 <- lm(mpg~horsepower,data=Auto)
```

```
automod <- augment(lm1, newdata = Auto) |>
```

```
  janitor::clean_names()
```

```
automod
```

```
# A tibble: 392 x 12
```

	rownames	mpg	cylinders	displacement	horsepower	weight	accel
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	
1	1	18	8	307	130	3504	
2	2	15	8	350	165	3693	
3	3	18	8	318	150	3436	
4	4	16	8	304	150	3433	
5	5	17	8	302	140	3449	
6	6	15	8	429	198	4341	
7	7	14	8	454	220	4354	
8	8	14	8	440	215	4312	
9	9	14	8	455	225	4425	
10	10	15	8	390	190	3850	

```
# i 382 more rows
```

```
# i 4 more variables: origin <dbl>, name <fct>, fitted <dbl>, re
```

```

autoboot <- automod |> specify(response = resid) |>
  generate(reps = 1000, type = "bootstrap") |>
  #ungroup() |>
  mutate(fitted = automod$fitted) |>
  mutate(new_mpg = fitted + resid) |>
  mutate(horsepower = automod$horsepower)

```

```
autoboot
```

```
# A tibble: 392,000 x 5
```

```
# Groups:   replicate [1,000]
```

	replicate	resid	fitted	new_mpg	horsepower
	<int>	<dbl>	<dbl>	<dbl>	<dbl>
1	1	-4.62	19.4	14.8	130
2	1	-2.05	13.9	11.8	165
3	1	0.113	16.3	16.4	150
4	1	-2.07	16.3	14.2	150
5	1	-5.26	17.8	12.6	140
6	1	-5.85	8.68	2.84	198
7	1	-2.57	5.21	2.64	220
8	1	1.07	6.00	7.07	215
9	1	-5.97	4.42	-1.54	225
10	1	-3.04	9.95	6.90	190

```
# i 391,990 more rows
```



```
slope_autoboot <- autoboot |>  
  group_by(replicate) |>  
  summarize(slope = coef(lm(new_mpg~horsepower))[2])
```

```
slope_autoboot
```

```
# A tibble: 1,000 x 2
```

	replicate	slope
	<int>	<dbl>
1	1	-0.170
2	2	-0.157
3	3	-0.160
4	4	-0.165
5	5	-0.150
6	6	-0.161
7	7	-0.169
8	8	-0.151
9	9	-0.166
10	10	-0.172

```
# i 990 more rows
```

```
percentile_ci <- quantile(slope_autoboot$slope, c(0.025, 0.975))
```

```
percentile_ci
```

2.5%	97.5%
-0.1688020	-0.1446102

Lab exercise 2

hake.csv contains abundance data for silver hake from tows in the 2015 NMFS spring bottom trawl survey.

- Produce 5,000 bootstrapped estimates for the mean abundance per tow based on case resampling (5000 samples).
- Compare the standard deviation of the bootstrapped estimates of the mean to the standard error of the mean from the original sample.
- Compute an approximate 95 percent confidence interval for the mean based on the bootstrap, assuming normality. Compare this to the interval based on percentiles of the bootstrap sampling distribution.
- BONUS* Plot how the bootstrap confidence interval for the mean changes with the number of bootstrap samples. (100, 500, 1000, 2000, 5000, 10000)

Validation Approach

Recall:

Split data into training set and a validation (or hold-out) set.

Fit model to training set, fitted model used to predict the responses for the observations in the validation set.

Resulting validation set error rate (e.g. MSE for quantitative response) is an estimate of the test error rate.


```
set.seed(1)
train <- sample(392,196)
lm.fit <- lm(mpg ~ horsepower, data=Auto, subset=train)
mean((Auto$mpg - predict(lm.fit, Auto))[-train]^2)
[1] 23.26601
```

```
lm.fit2 <- lm(mpg ~ poly(horsepower, 2), data=Auto, subset=train)
mean((Auto$mpg - predict(lm.fit2, Auto))[-train]^2)
[1] 18.71646
```

```
lm.fit3 <- lm(mpg ~ poly(horsepower, 3), data=Auto, subset=train)
mean((Auto$mpg - predict(lm.fit3, Auto))[-train]^2)
[1] 18.79401
```

Wages (from ISLR)

```
ggplot(Wage, aes(y = wage, x = age)) +  
  geom_point(col = "slateblue", alpha = 0.1, size = 4) +  
  theme_minimal() #+
```



k-Fold Cross-Validation

Use `cv.glm()` with argument `K=k` to perform k-fold cross-validation.

```
set.seed(17)
cv.error.10 <- matrix(0,nrow=10,ncol=10)
for (isim in 1:10) {
  cv.error <- map(1:10,
    ~cv.glm(Auto,
      glm(mpg~poly(horsepower,.x),data = Auto),
      K=10)$delta[1])
  cv.error.10[isim,] <- as.numeric(cv.error)
}
```

`cv.error.10`

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	
[1,]	24.27207	19.26909	19.34805	19.29496	19.03198	18.89781	19.12061	19
[2,]	24.32221	19.20485	19.20345	19.19061	18.81104	18.93199	18.64288	18
[3,]	24.40176	19.28566	19.47768	19.45928	19.11895	18.96636	19.14254	19
[4,]	24.34410	19.21850	19.23349	19.29482	19.20885	18.89513	18.83750	19
[5,]	24.12350	19.15993	19.30270	19.43933	19.14562	18.77230	18.95511	18
[6,]	24.26450	19.20361	19.21035	19.50279	19.03784	19.00893	18.80578	20
[7,]	24.21276	19.32266	19.25201	19.57620	18.94538	18.98961	19.36863	18
[8,]	24.23926	19.19939	19.36453	19.18138	18.98821	19.20072	18.82207	19
[9,]	24.25723	19.19239	19.43090	19.64706	19.03607	18.83338	18.74457	18
[10,]	24.06964	19.23239	19.41729	19.54851	19.02514	18.90808	18.94559	19
	[,9]	[,10]						

Lab exercise 3, k-fold cross validation

Using the Wage data set, evaluate the predictive ability of models for wages.

- a. Define a unique random number seed. Use 10-fold cross validation to estimate the test error rate for models fitting a polynomial of age of order 2, 3, 4, 5, and 6.
- b. Conduct the validation 20 times for each polynomial. Plot the distribution (use boxplots) for the validation test error rate as a function of the degree of polynomial. Based on the results, what order polynomial would you use?
- c. Use 5-fold cross-validation to compare the performance of models that include combinations of:
 - a polynomial of age
 - education level
 - race
 - industry
- e. What model would you choose based on the test error rates?
- f. *BONUS* How does the model chosen by 5-fold CV compare to that from using AIC as a model selection tool?

HINTS

Write down the steps you need to take to perform the calculations.

Make use of existing code from earlier in the lab to help.

Lab schedule

1/22: Introduction to R and R Studio, working with data

1/29: Visualizing data

2/05: Probability, linear modeling

2/12: Data wrangling, model summaries

2/20: Simulation, Resampling

2/26: Iteration

3/19: Creating functions, debugging

4/02: working with Spatial data

4/09: Flex: automated reporting & Quarto